

AI Decision Exposure Brief

Agent Audit Kit (AAK)

February 23, 2026

Audience: Regional bank CTO/CISO, risk, internal audit, and technology leadership.

Executive Summary

Most institutions now have AI-assisted workflows in operations, support, underwriting, fraud review, or internal analysis. The operational risk is not only model quality. The larger risk is *decision exposure*: if an AI-linked decision is challenged, can your team reconstruct what happened in a defensible way?

Many teams can produce logs. Fewer can produce a replayable, tamper-evident decision lineage that stands up under incident response, legal review, or regulator scrutiny.

The Core Gap

Logging answers: “what events were recorded?”

Defensible lineage answers: “can we reproduce the exact path of prompts, tool calls, and outputs, and prove artifacts were not silently altered?”

Including multi-turn psychological convergence patterns (authority + urgency + social proof) that traditional logs miss.

In practice, the gap appears when teams must answer:

- Why did this automated action occur?
- Which external tools and data sources influenced the outcome?
- What evidence is missing today if this decision is disputed?
- Could we reproduce this review path tomorrow with confidence?

What We Do

We provide a fixed-scope **AI Decision Exposure Review** for one selected workflow.

- Map one real workflow where AI influences a material decision.
- Capture and reconstruct decision lineage using deterministic replay methods.
- Identify irreversible actions, external side effects, and evidence gaps.

- Produce a concise, executive-readable exposure report with concrete next steps.

Reference scenario baseline: Validated in deterministic synthetic CEO-fraud and Command Authority Confusion scenarios matching financial services attack patterns. Live workflow validation is conducted per-client during engagement.

What We Do Not Claim

- We do **not** certify compliance.
- We do **not** claim model correctness or safety.
- We do **not** infer intent.
- We do **not** replace legal, compliance, or internal audit functions.

Engagement Structure (Fixed Scope)

Duration: 1–2 weeks

Format: Document review + workflow walk-through + technical evidence pass

Primary Output: AI Decision Exposure Report

Typical report sections:

- Workflow boundaries (where AI influences decisions)
- Data ingress and egress points
- High-impact decision points and irreversible actions
- Current replay and evidence-readiness score (practical, not theoretical)
- Deterministic narrative report + Mermaid convergence timeline + golden-run comparison
- Prioritized remediation actions (30/60/90 day)

Example Decision Convergence Timeline

The delivered exposure report includes a decision-timeline table in this format:

Time	Authority	Temporal	Social	Conv.	Risk	Action
10:30	70%	55%	—	0.15	YELLOW	Enhanced review
11:00	61.6%	75%	—	0.20	RED	Blocked
11:15	65%	66.3%	68%	0.30	CRITICAL	SOC alert

Investment

Pilot range: \$2,500–\$5,000 per workflow review, fixed fee.

This price point is designed for rapid internal approval and practical decision-making, without committing to a platform purchase.

Decision Question for Leadership

If one AI-influenced decision is challenged next quarter, can we prove what happened with evidence that is complete, replayable, and defensible?

If the answer is uncertain, this review is intended to close that gap quickly and concretely.